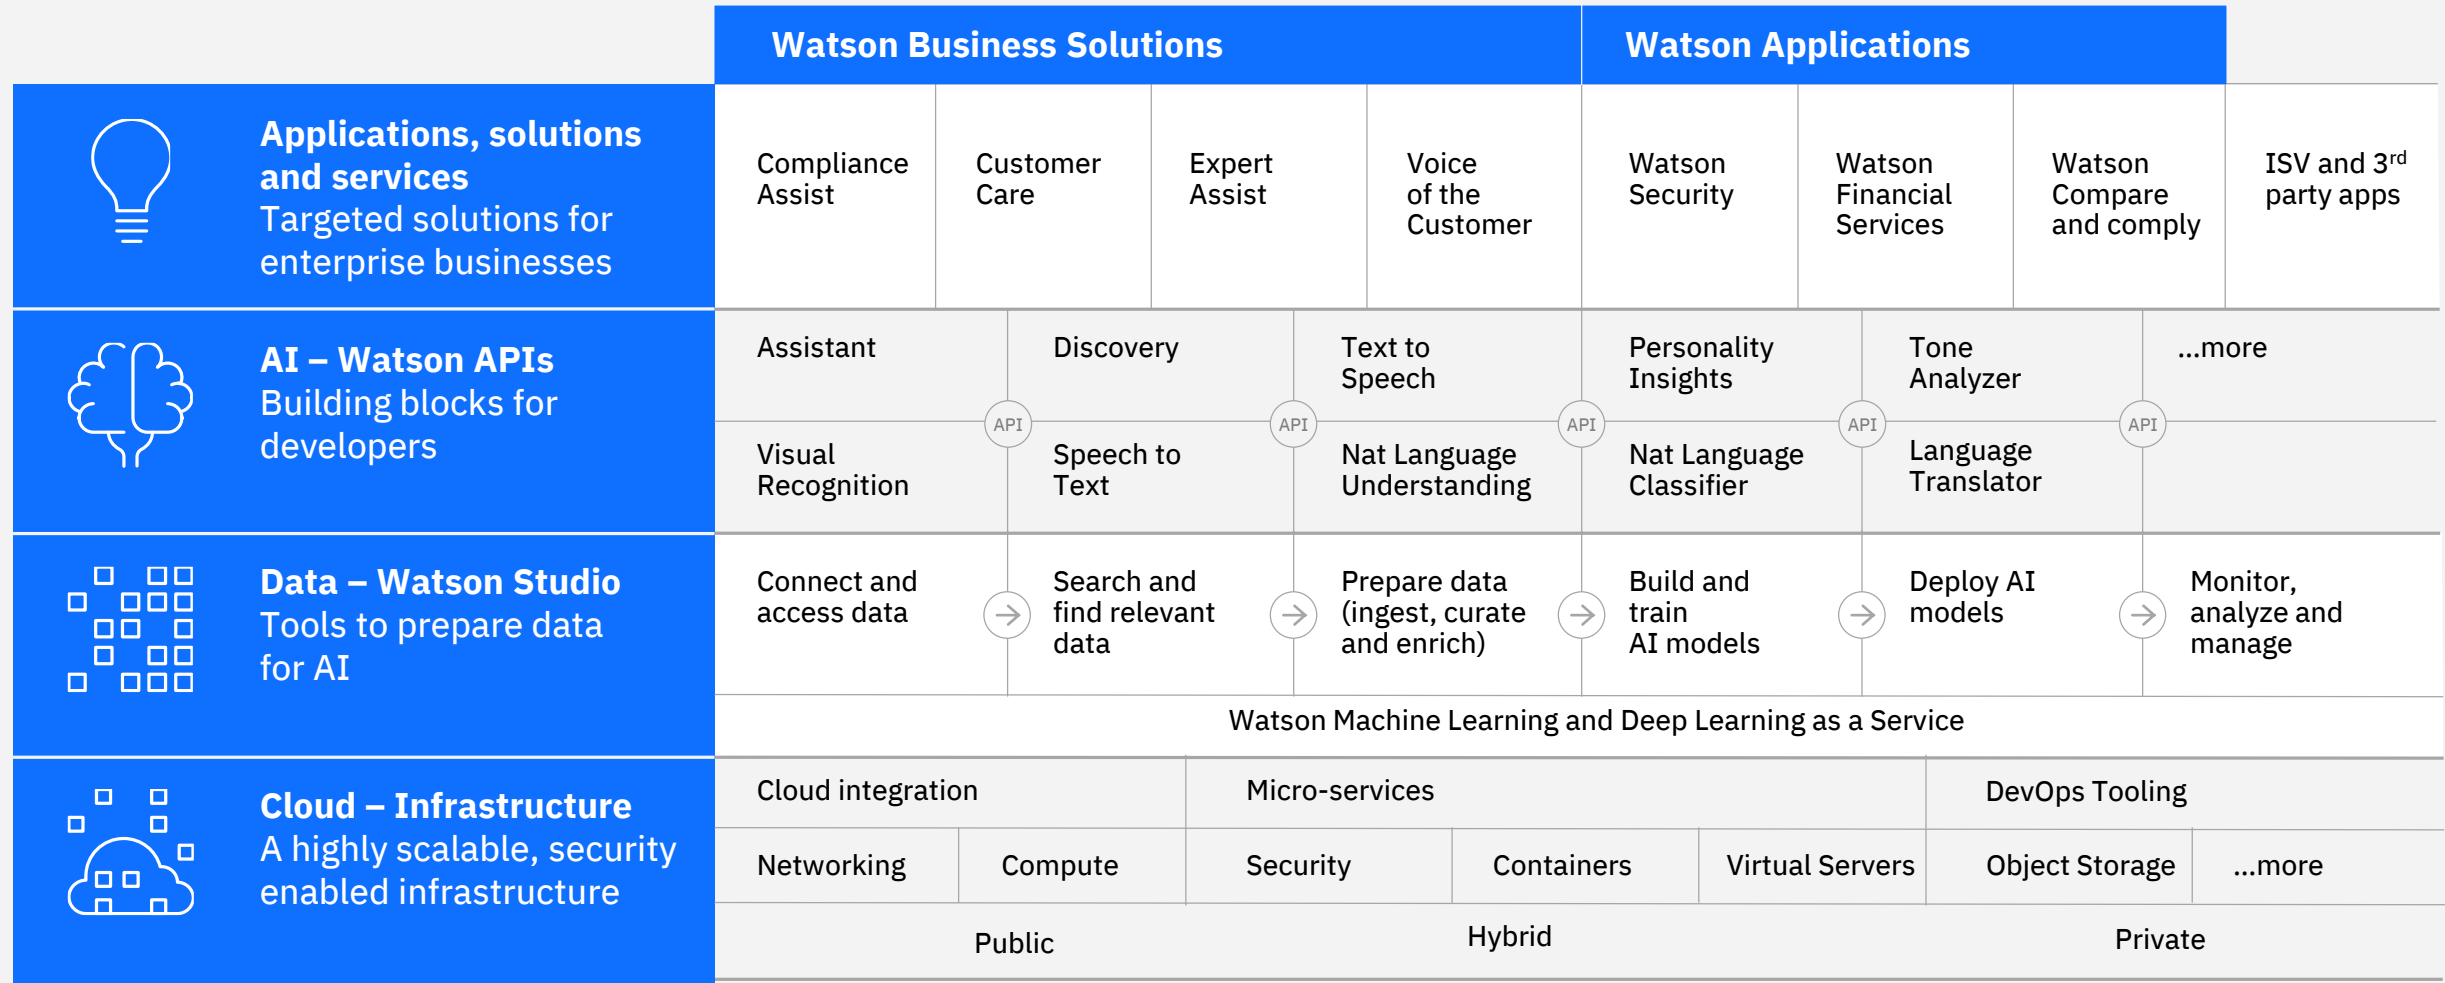


Agenda:


1. Overview of IBM Offerings in text analytics - Intro (5 min)
2. Service details: (60 min)
 - a) Watson Discover
 - b) Watson NLU
 - c) Watson Tone Analyzer
 - d) Watson NLC
 - c) Watson Knowledge Studio
3. Demo - STB dataset (1 hr)
4. Hands on using STB dataset (1hr)
5. Q&A (15 min) & Wrap up

Watson is AI for smarter business




What are Watson APIs?

- **A collection of cloud services for developing Enterprise AI solutions**
Protects your insights
- **Services can be trained using custom models**
Learn more with less data
- **Each service has a RESTful API**
Usable from virtually any language/platform
- **Provisioned at IBM Cloud**
- <http://ibm.com/watson>
Demos, documentation, API reference, sample apps, SDKs
- <https://developer.ibm.com/code/patterns/>
Comprehensive AI code patterns




Four capabilities distinguish Watson from traditional programmed computing systems



Understands

imagery, language and other
unstructured data humans




Reasons

forms hypotheses, infers
and extracts ideas



Learns

With every interaction to
sharpen expertise



Interacts

With humans through
seeing, talking and hearing

Language



Assistant



Language
Translator



Personality
Insights



Tone
Analyzer



Natural Language
Understanding

Speech



Speech to
Text



Text to
Speech

Vision



Visual
Recognition

Data Insights



Discovery

Watson APIs



Watson Assistant



Language Translator



Speech to Text




Tone Analyzer



Natural Language Understanding



Natural Language Classifier



Personality Insights



Visual Recognition



Discovery

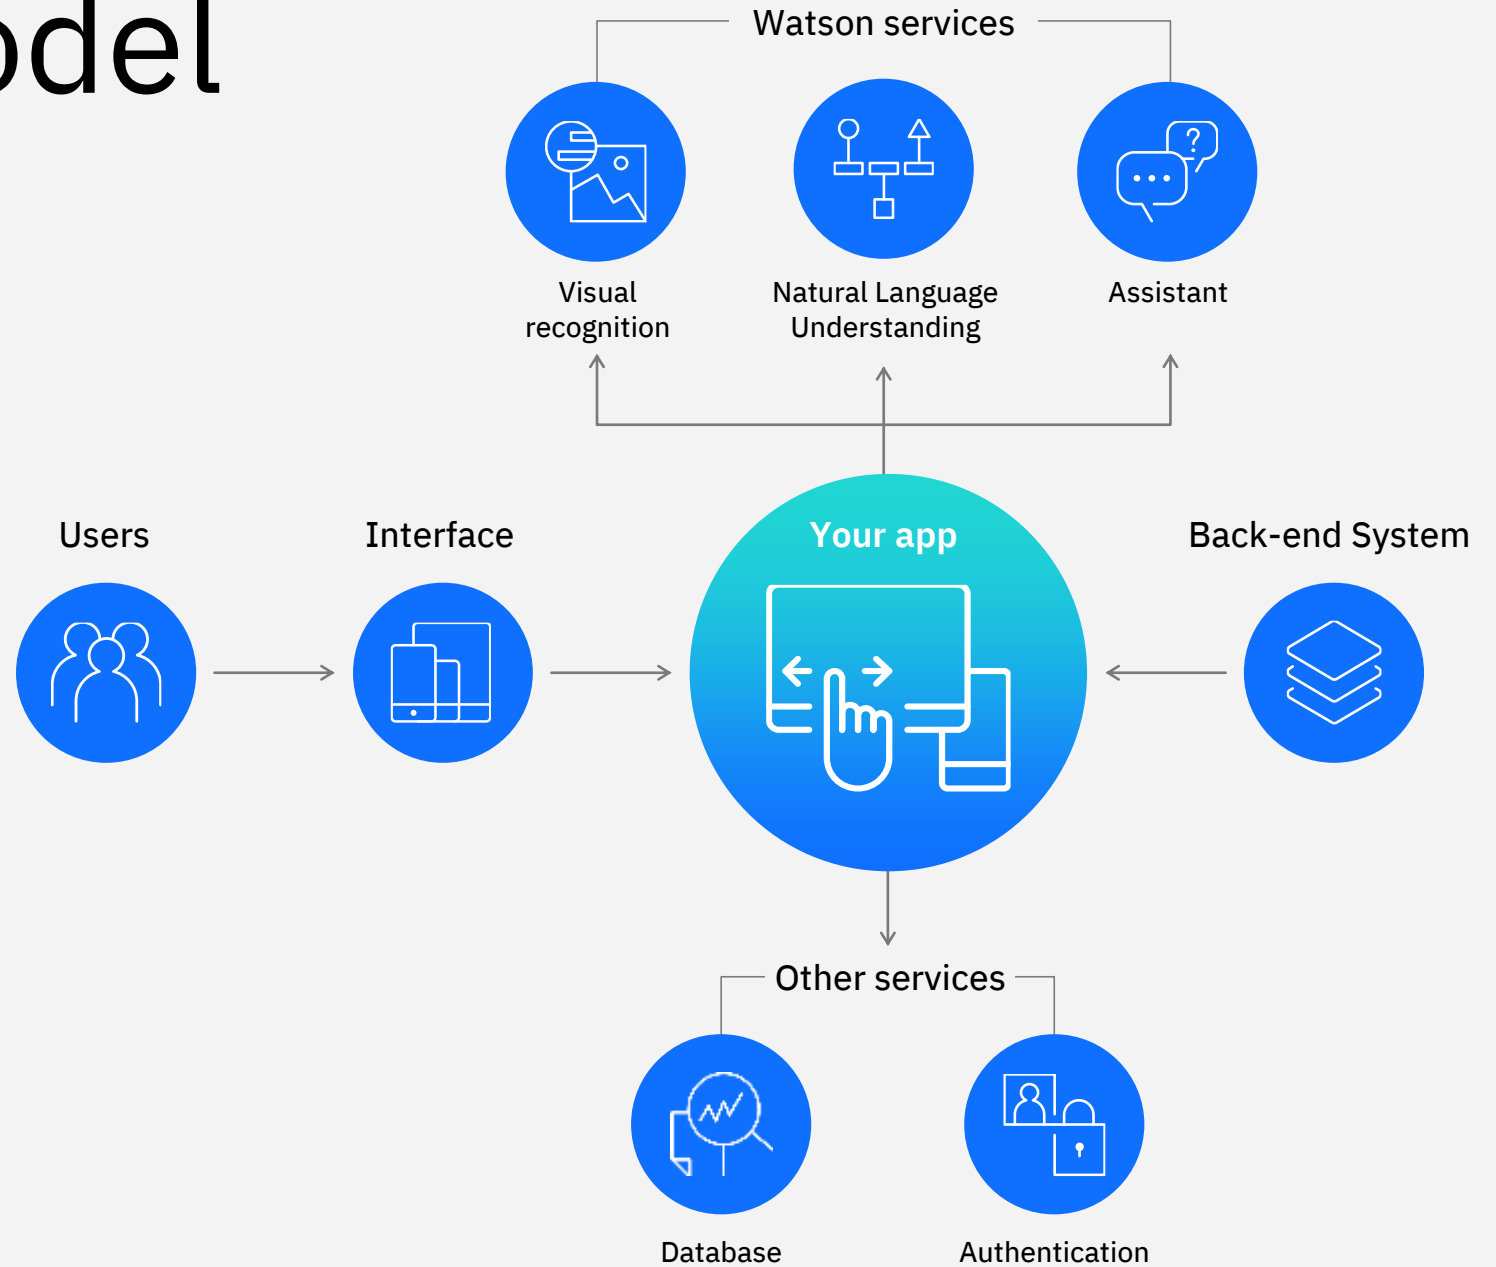


Text to Speech

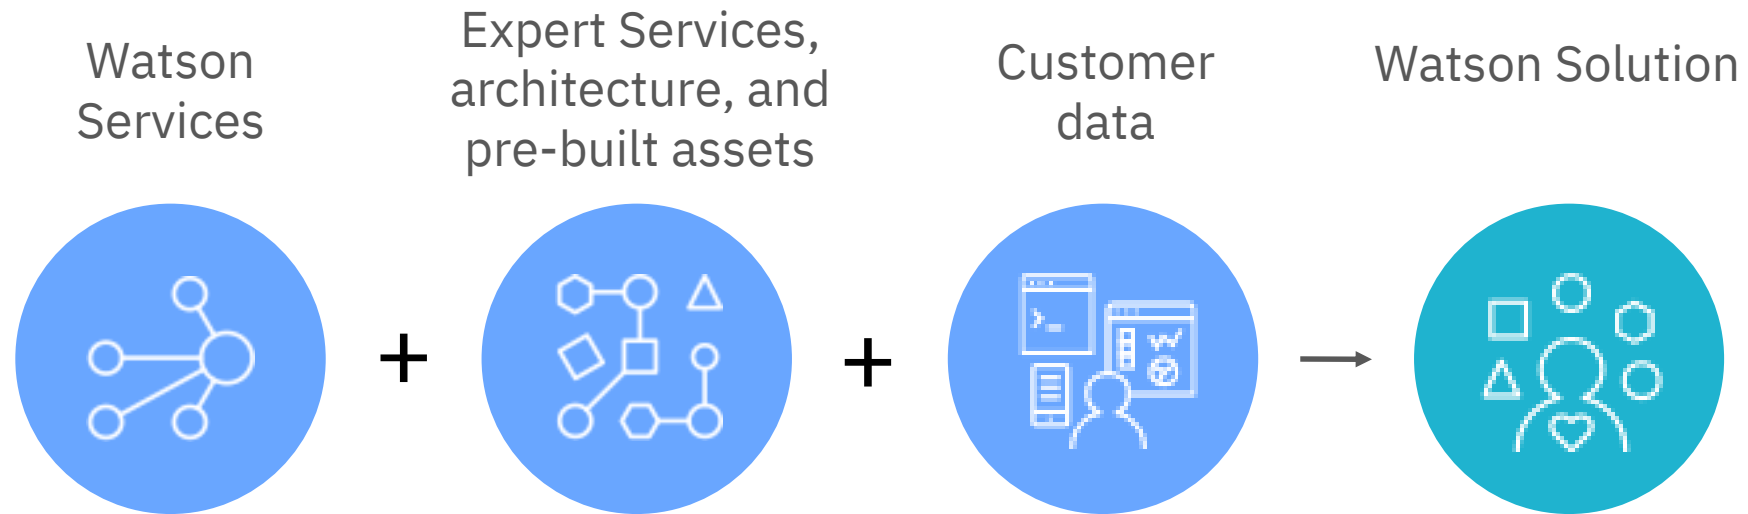
Execution model

A main application

- Hosted on IBM Cloud
- Running on servers that you manage



Watson Solutions combine Watson Services, Watson APIs, and pre-built assets to quickly prove value and deliver foundational applications



IBM Watson NLU



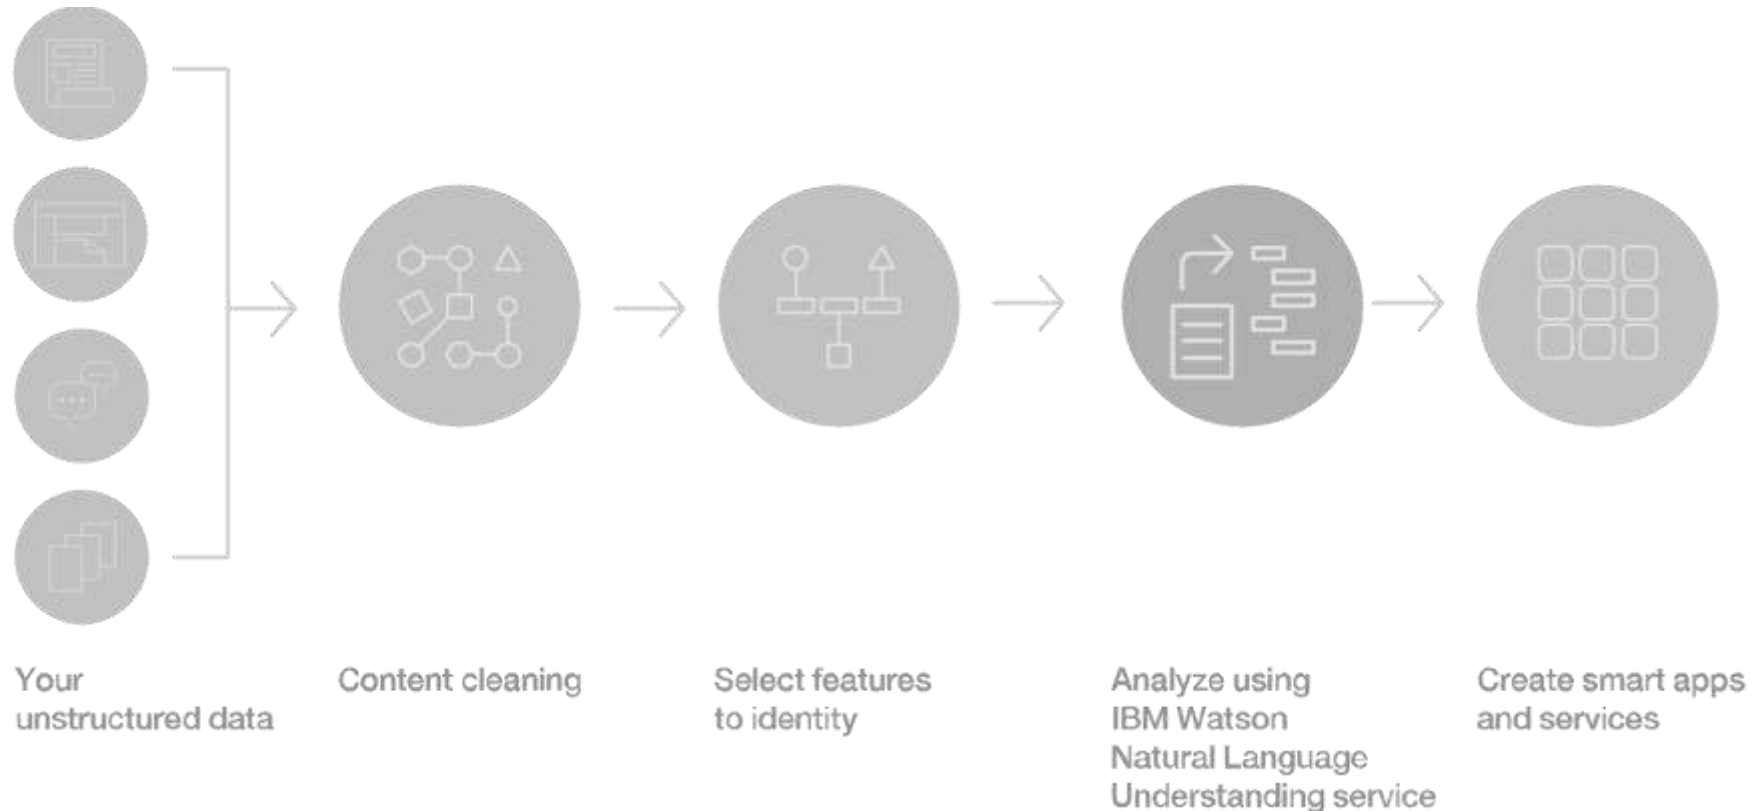
Natural language processing for advanced text analysis

A sophisticated suite of natural language processing capabilities to analyze text and extract meta-data from content such as concepts, entities, keywords, categories, sentiment, emotion, relations, semantic roles, with options for customization to specific industries and domains.

How does it work?

How it works

IBM Watson Natural Language Understanding makes value-driven decisions easy by giving you the full story behind all of your data.



Highlights

Comprehensive NLP stack: A rich set of enrichments to extract metadata from unstructured text for deep analysis in various use cases

Broad Language Coverage: Depending on the feature, NLU can understand text in thirteen different languages

Customization: Industry and domain-specific annotation possible with custom models developed using Watson Knowledge Studio resulting in higher accuracy broader uses

Improved API: API Signature is consistent with all Watson APIs enabling quicker development

Scalability: Cloud-based service with the infrastructure to support millions of NLU enrichments per month

**What are the
main
features?**

Features

Keywords: Determine the most important keywords in your content.

Concepts: Identify general concepts in your content.

Categories: Categorize your content into a hierarchical 5-level taxonomy.

Entities: Detect important people, places, geopolitical entities and other types of entities.

Sentiment: Determine whether your content conveys positive or negative sentiment.

Features (cont.)

Emotion: Detect emotions such as anger, disgust, fear, joy or sadness that are conveyed by your content.

Relations: Identify relationships between entities in your content.

Semantic roles: Identify the subjects of actions, and the objects that they act upon.

Metadata: Get author information, publication date and the title of your text or HTML content.

Custom models: Use IBM Watson Knowledge Studio to customize text extraction.

**What are
the use
cases?**

Applications

Content Recommendation: How can I recommend content that readers might like?

Advertising Optimization: Where and how should I place my ads to the right audience?

Audience Segmentation: How can I segment customers to tailor strategies for different audiences?


Voice of Customer Analysis: How can I address customer feedback and improve customer experience?

Data Mining: How can I search a corpus and discover insights in an efficient way?

Use Case

Content Recommendation

Recommend similar content to what a user is interested in viewing or experiencing



Use Case

Advertisizing Optimization

Ensure proper placement of advertisements based on page content and viewer patterns



Use Case

Audience Segmentation

Identify key customer groups for market research purposes and to personalize campaigns towards different segments



Use Case

Voice-of-Customer Analysis

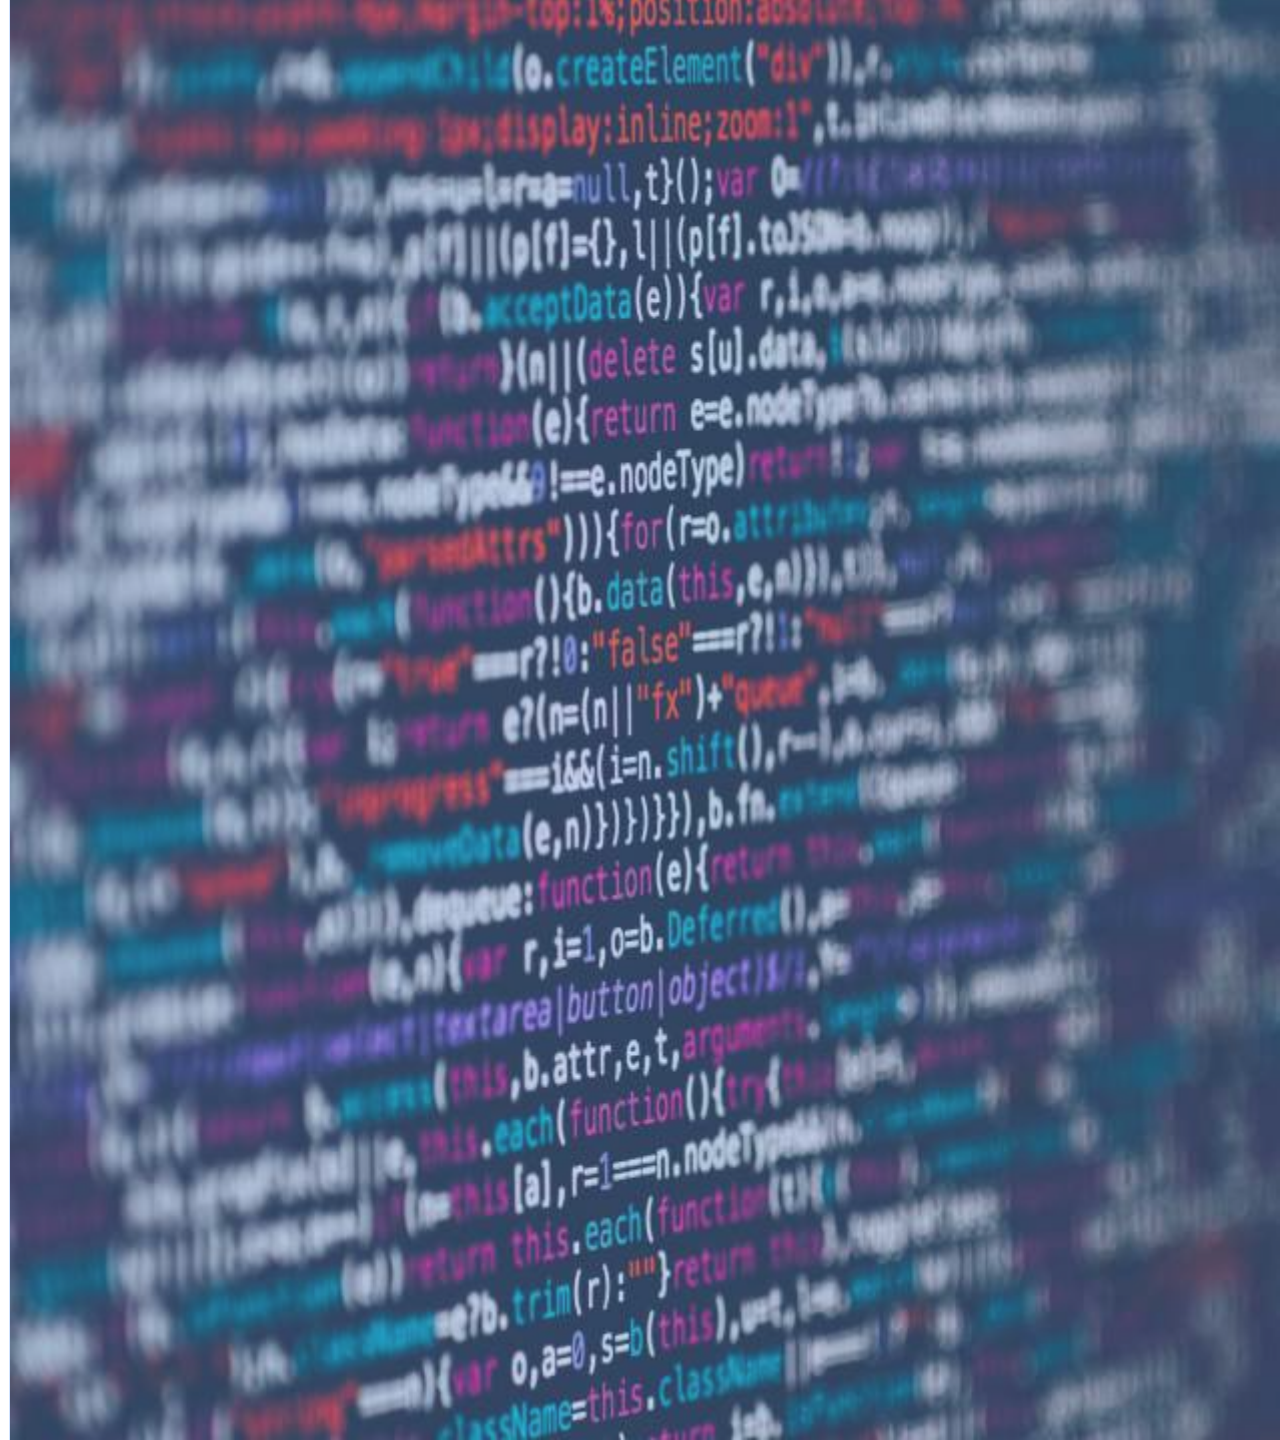
Spot trends in customer feedback to identify business opportunities, address concerns, reduce churn, and drive revenue



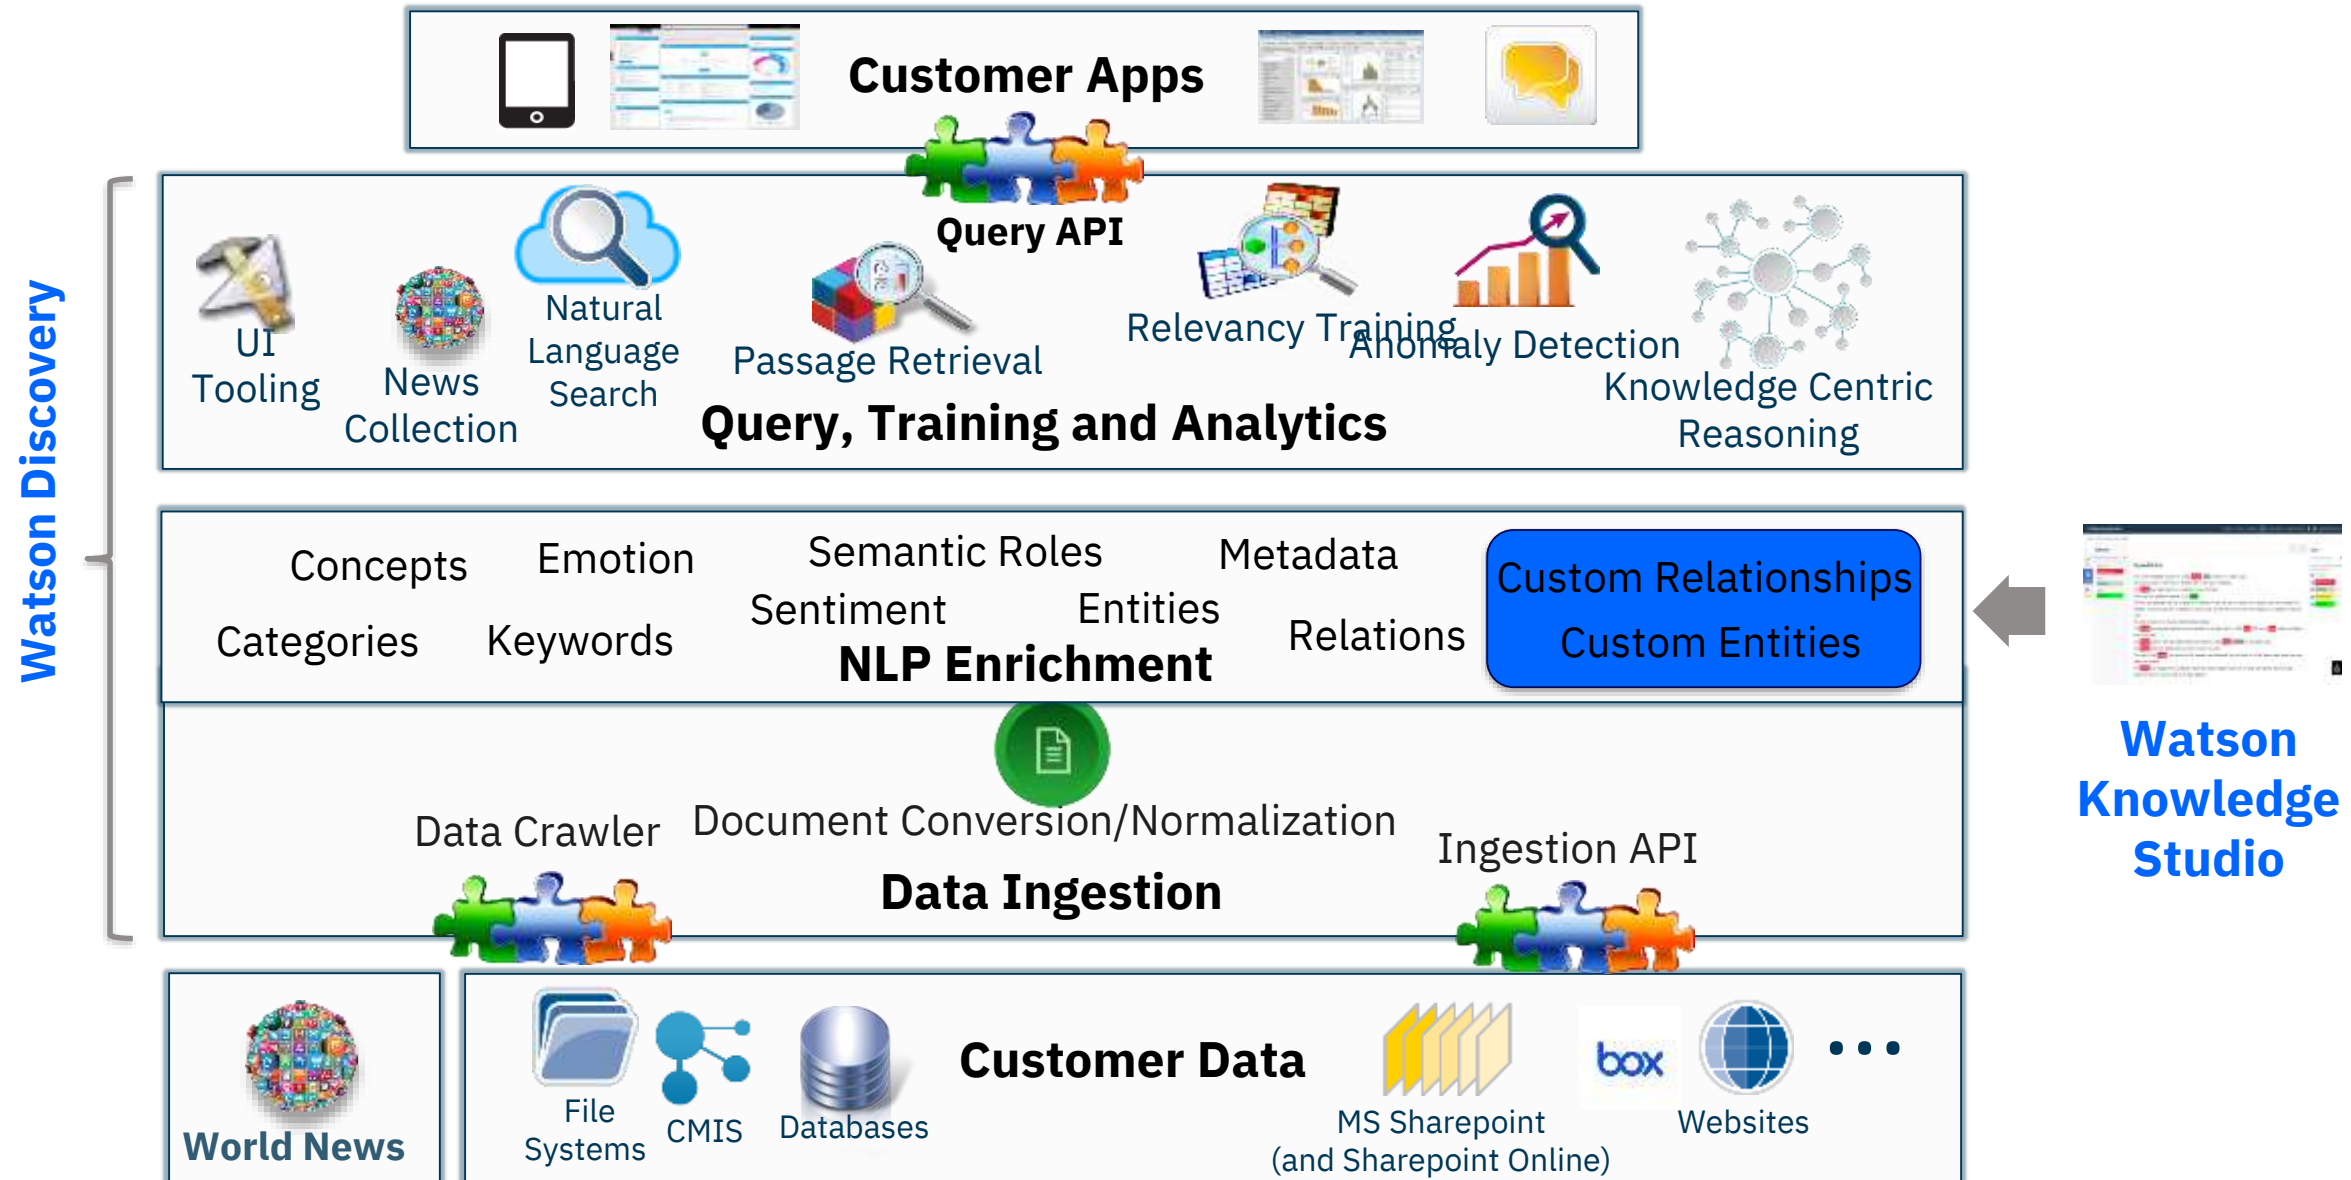
Use Case

Data Mining


Search through large content repositories and get answers that are relevant and quick



Watson Discovery Components



Watson Discovery : From Data to Insights - Beyond Enterprise Search



Documents and Data



"What phones are available that run Android and have good battery life under \$400?"

"What plans include international roaming?"

"What are the most frequently mentioned problems with this phone model?"

"What category of questions are customers the most upset about?"

Smart Document Understanding

Teach Discovery how to interpret document structure and components

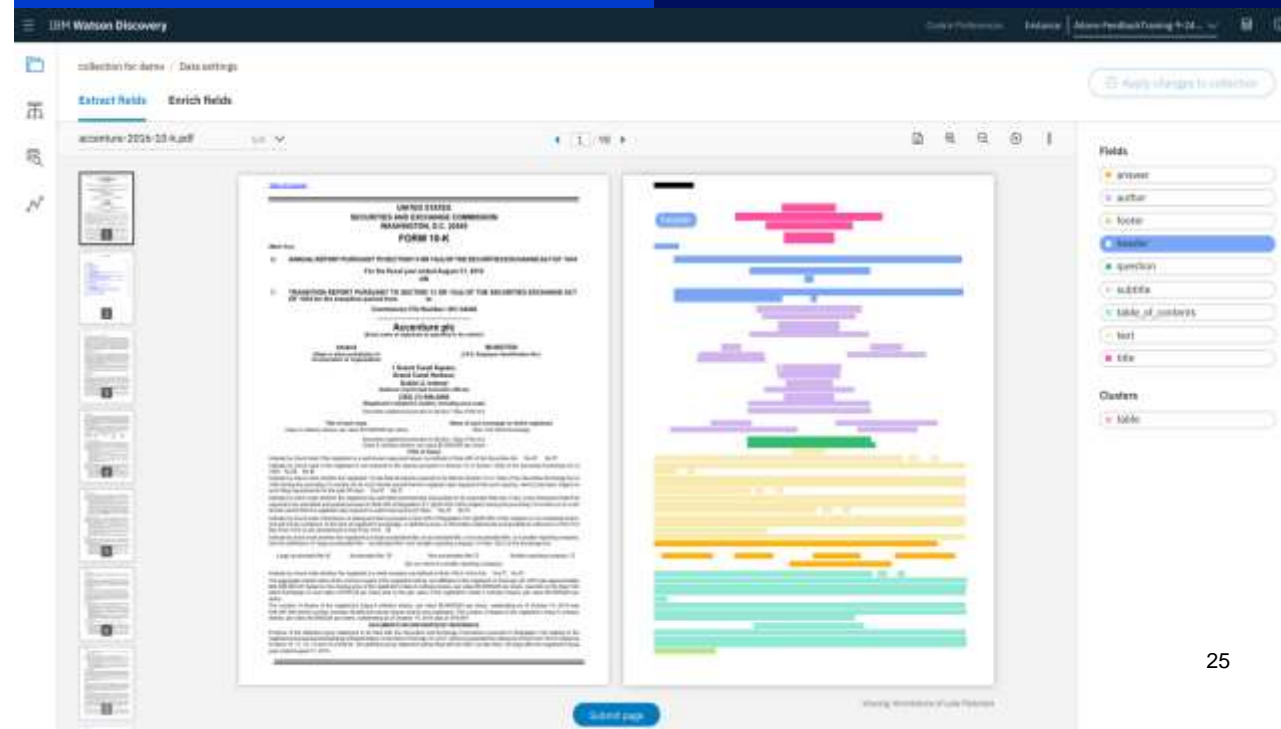
- Discover deeper insights by using extracted content (such as table header or abstract) in queries, to segment documents, ignore irrelevant parts and more
- Even identifies text in scanned business documents (OCR)
- Document formats supported: pdf, word, excel, powerpoint, html, png, jpeg, tiff and json
- Train Discovery to understand complex and domain specific tables

Rapidly Train

- Train AI model with point and click
- Label content in a small set of sample documents, including semantics like title, abstract, author or your own custom labels.
- Train on table semantics, header, footer, columns, more.
- Tests show 99% accuracy in as little as 200-300 labeled pages. Homogenous data sets in even less trained pages.

With labeled document sections you can :

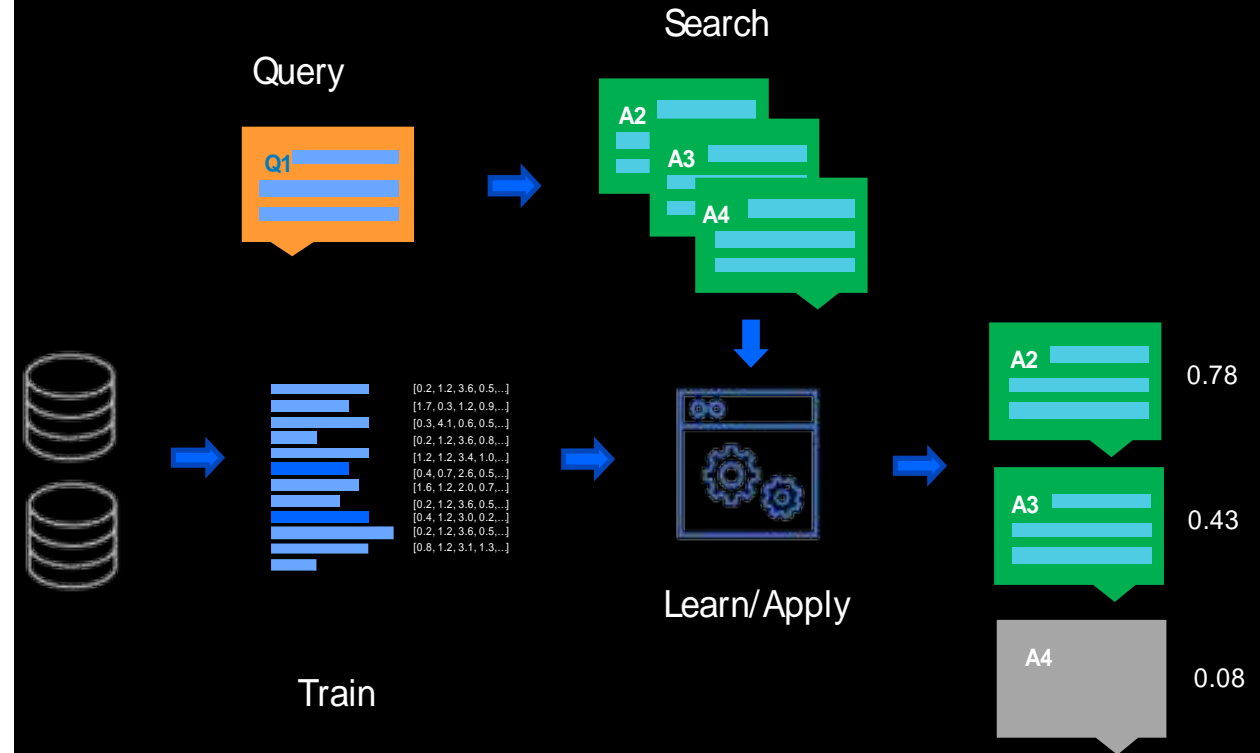
- *Not index certain content*
- *Merge content*
- *Split content into separate documents*
- *Identify Q&A Pairs*
- *Perform targeted queries on specific content*



Relevancy Training

Teach Discovery how to find relevant results with representative examples

- Increase overall relevancy of results across a large set of answer documents without training for specific content
- Scales to learn from general, corpus-wide signals that indicate relevance rather than query or document specific signals



Connectivity

Tooling based connectivity framework for easy loading of data

Available In Q2

- Box*
- SharePoint (Online and On-Prem)*
- Salesforce
- Web Crawl

Available in a July fixpack

- Fileshare*
- Database (ODBC)

* Includes support for Doc Level Security



Query confidence prediction


Discovery uses pre-trained models to estimate the relevance of a result and return a normalized confidence score

Allows for **setting thresholds** to improve overall relevance

Can be used in combination with Watson Assistant to determine when to use long-tail results


Planned Improvements:

- **More accurate** scoring of results using semantic and statistical features



Custom Model Enrichments

- Watson Knowledge Studio now included as part of Bundle/Add-On
- Custom Model enrichment support now available on WD on ICP4D



Watson Assistant Search Skill

- **Improve the coverage of your assistant** by calling out to a “search skill” from any dialog node
- **Search results are formatted automatically** by search skill
- **Ingestion and Sync to common knowledge base and CMS products**, like Salesforce Knowledge, Box, Sharepoint, and others...

Skills / DemoSkill

Configure Search Response

Search skill results will be surfaced to end users as a card. Map your data schema from Discovery to the title, body, and URL fields below to define what results will be surfaced to end users in the card results.

Title (optional)

title Example Web Page Title X

Body (optional)

body Example Card Body Text For example, card 1... X

URL (optional)

url Example This card points to Search Information... X

Define the text your search skill will display to the end user

Message No results found Connectivity issue

I searched my knowledge base and found this information which might be useful:

Try it out Clear X

Preview how search skill results will be surfaced

is lost luggage covered

I searched my knowledge base and found this information which might be useful:

Lost Luggage

The benefit amounts for jewelry, watches, cameras, video recorders, and other electronic equipment are part of and not in addition to

["Benefit Information What is the Lost Lugga...]

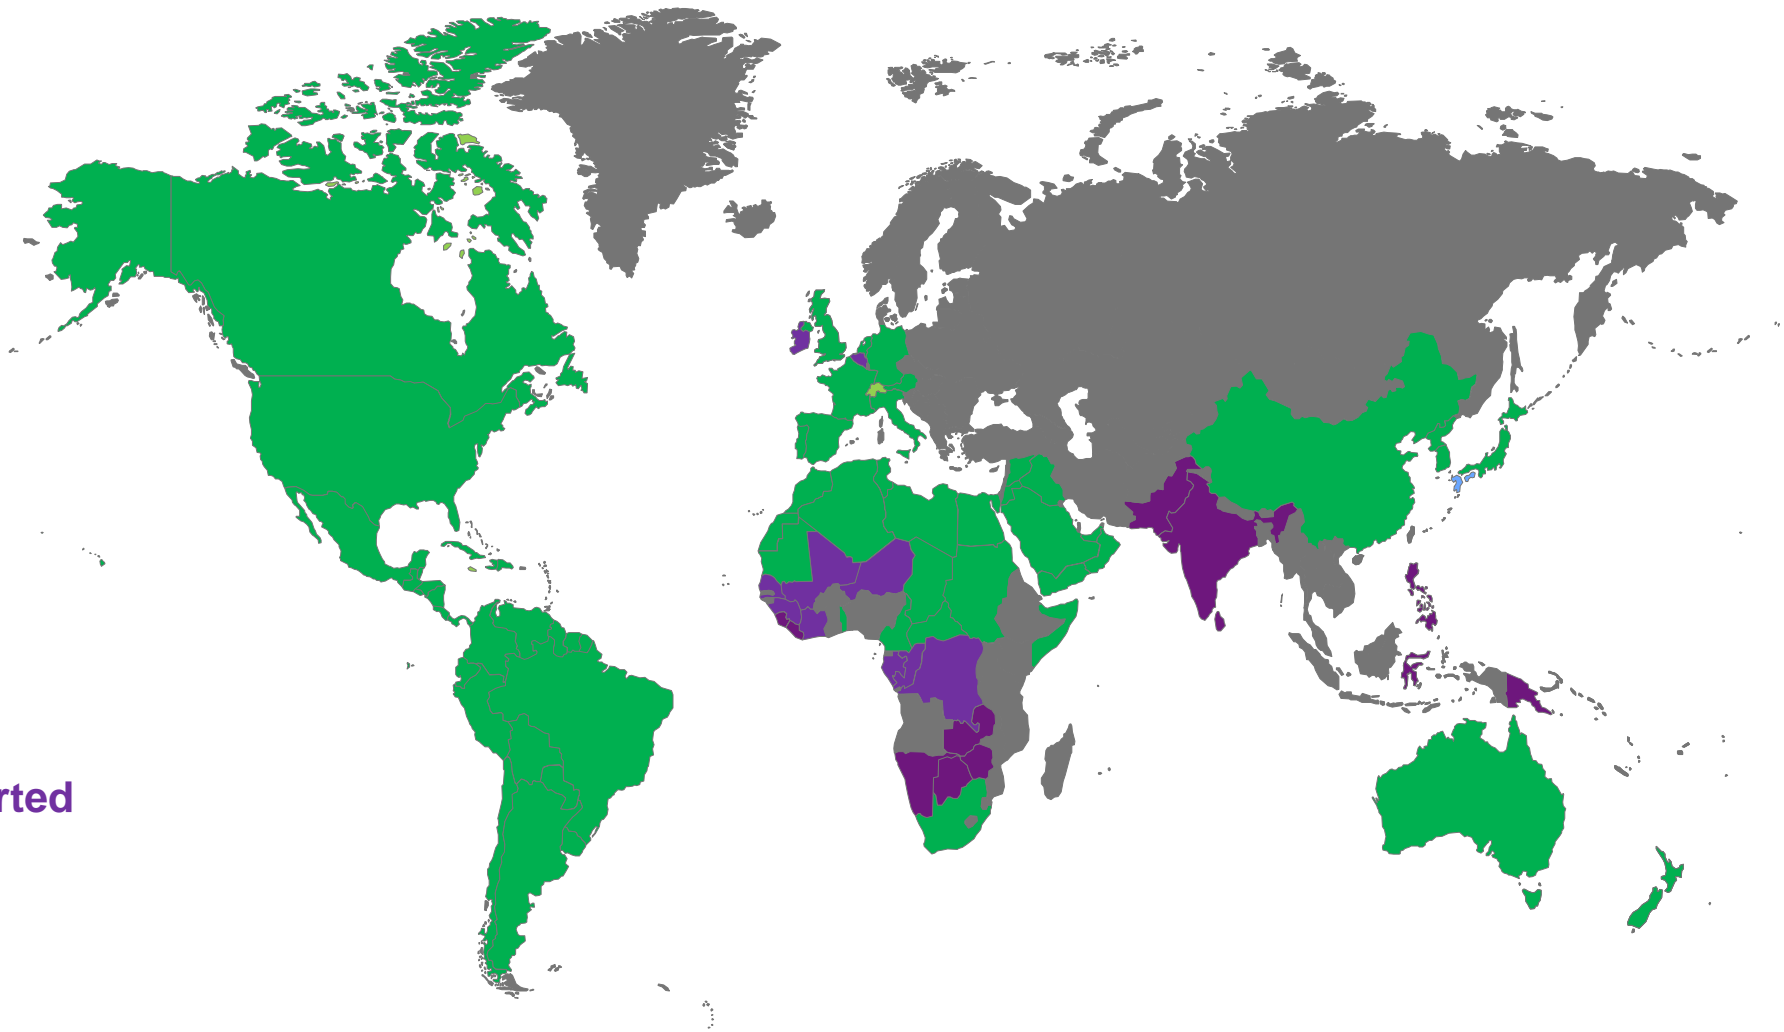
["Travel and Emergency Assistance Services"]

• Lost Luggage Locator Service can help you through the common carrier's claim procedures or can arrange shipment of


["What are Travel and Emergency Assistance ..."]

Language Support (Same on Public and IBM Cloud Pak for Data)

- Arabic
- Japanese
- Dutch
- Chinese Simplified
- English
- Portuguese (BR)
- Italian
- German
- Korean
- Spanish
- French
- Some Languages Supported**



WA + Discovery for all types of Qs



Watson Assistant



Watson uses reasoning strategies that focus on the language and context of the question to answer simple questions.

Watson Discovery




Watson uses reasoning strategies that focus on identifying the most appropriate answer for complex questions using the context of the entire corpus.

Watson Knowledge Studio



Teach Watson the language of your domain.



Complication

Ambiguity of natural language makes it complex for people to train machines to read it.



Isolation

Isolated development environments make it challenging for domain experts and developers to collaborate.



Application

Custom natural language models are difficult to actually apply to data, requiring further technical work and complex architectures



Teach by Example

Empower domain subject matter experts to teach Watson the language of your industry or organization without requiring deep technical skills or coding.



Engage your Experts

Bring your experts together into a collaborative online environment to teach Watson how to understand the linguistic nuances of your domain.



Use Everywhere

Apply what you've taught Watson in multiple applications by using models from Watson Knowledge Studio in Watson Discovery, Watson Natural Language Understanding and Watson Explorer.

Demo